Summary

Let's look at where we are now

1. We have covered a lot of interesting topics in regularization
2. We haven't covered the regularization methods such as dropout & batch-normalisation, but they will be covered as we move forward
3. The next few sections will be more hands-on, and we will get to start working with PyTorch and CNNs
4. The next contest will cover all of the following concepts

Network Architectures
- Number of layers
- Number of Neurons

Algorithms
- Vanilla/momentum/NAG
- AdaGrad
- RMSProp
- Adam

Activation Functions
- tanh (RNNs)
- ReLU (CNNs, DNNs)
- leaky ReLU (CNNs)

Next Contest

Strategies
- Batch
- Mini-Batch (32, 64, 128)
- Stochastic
- Learning Rate schedule

Regularization
- L2
- Early Stopping
- Dataset Augmentation
- Drop-out
- Batch Normalization

Initialization Methods
- Xavier
- He

* red text indicates concepts not covered as of yet.